

<https://iaarbook.github.io/img/computer-vision.jpg>

RUHR-UNIVERSITÄT BOCHUM

BASICS ON (CONVOLUTIONAL) NEURAL NETWORKS

How to teach a computer numbers

Computational Physics II - 25/07/2022 - Jurek Völp, jurek.voelp@ruhr-uni-bochum.de

Outline

- Basics of neural networks

How a machine learns

- Convolutional Neural Network

How a machine learns to read numbers


- Summary

Basics of neural networks

How a machine learns

Neuron

- Basic part in a neural network
- Mathematical representation:
 - $\hat{y} = g(\sum_i x_i w_i)$
 - $g(z) \hat{=}$ *Activation Function*
 - Sigmoid
 - Hyperbolic Tangent
 - Relu
 - $x_i \hat{=}$ *Inputs*
 - $w_i \hat{=}$ *Weights*
 - How important is a certain input?



<http://introtodeeplearning.com>, MIT 6.S191

Neuron – Implementation

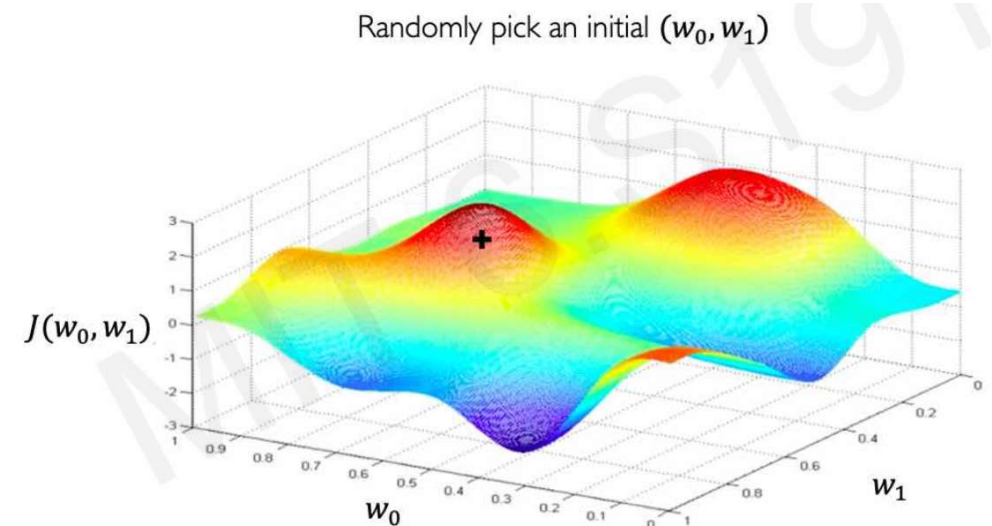
```
1 import numpy as np
2 class Neuron:
3     def __init__(self, numberOfInputs, actFunc = "sigmoid"):
4         # generate random weights
5         self.weights = 2 * np.random.random(numberOfInputs) - 1
6         self.bias = 2 * np.random.random() - 1
7         # select the activation function
8         if type(actFunc) == str:
9             if actFunc == "sigmoid":
10                 self.activation_func = lambda z: 1 / (1 + np.exp(-z))
11                 self.gradActivation_func = lambda z: self.activation_func(z) * (1 - self.activation_func(z))
12             elif actFunc == "hyperbolicTangent":
13                 self.activation_func = lambda z: (np.exp(z) - np.exp(-z)) / (np.exp(z) + np.exp(-z))
14                 self.gradActivation_func = lambda z: 1 - self.activation_func(z)**2
15             elif actFunc == "relu":
16                 self.activation_func = lambda z: np.maximum(0, z)
17                 self.gradActivation_func = lambda z: np.heaviside(z,0)
18             elif callable(actFunc):
19                 self.activation_func = actFunc
20             else:
21                 raise Exception("No valid activation function given.")
22
23     def __call__(self, input):
24         if np.size(input, axis=1) == self.weights.size:
25             return self.activation_func(self.bias + np.dot(self.weights, input.T))
26         else:
27             raise Exception("Input is not in the right size.")
```

Neuron – Implementation

```
1 import numpy as np
2 class Neuron:
3     def __init__(self, numberOfInputs, actFunc = "sigmoid"):
4         # generate random weights
5         self.weights = 2 * np.random.random(numberOfInputs) - 1
6         self.bias = 2 * np.random.random() - 1
7         # select the activation function
8         if type(actFunc) == str:
9             if actFunc == "sigmoid":
10                 self.activation_func = lambda z: 1 / (1 + np.exp(-z))
11                 self.gradActivation_func = lambda z: self.activation_func(z) * (1 - self.activation_func(z))
12             elif actFunc == "hyperbolic tangent":
13                 self.activation_func = lambda z: (np.exp(z) - np.exp(-z)) / (np.exp(z) + np.exp(-z))
14                 self.gradActivation_func = lambda z: 1 - self.activation_func(z)**2
15             elif actFunc == "relu":
16                 self.activation_func = lambda z: np.maximum(0, z)
17                 self.gradActivation_func = lambda z: np.heaviside(z,0)
18             elif callable(actFunc):
19                 self.activation_func = actFunc
20             else:
21                 raise Exception("No valid activation function given.")
22
23     def __call__(self, input):
24         if np.size(input, axis=1) == self.weights.size:
25             return self.activation_func(self.bias + np.dot(self.weights, input.T))
26         else:
27             raise Exception("Input is not in the right size.")
```

How to learn/train


- The cost/loss function quantifies how incorrect the prediction is!
 - Function of the weights
 - Examples:
 - Binary Cross Entropy Loss – between 0 and 1
 - Mean Squared Error Loss – any real number
- **Aim of the training: Optimize/Reduce the cost function!**
- Gradient descent:
 - random initial values



<http://introtodeeplearning.com>, MIT 6.S191

How to learn/train

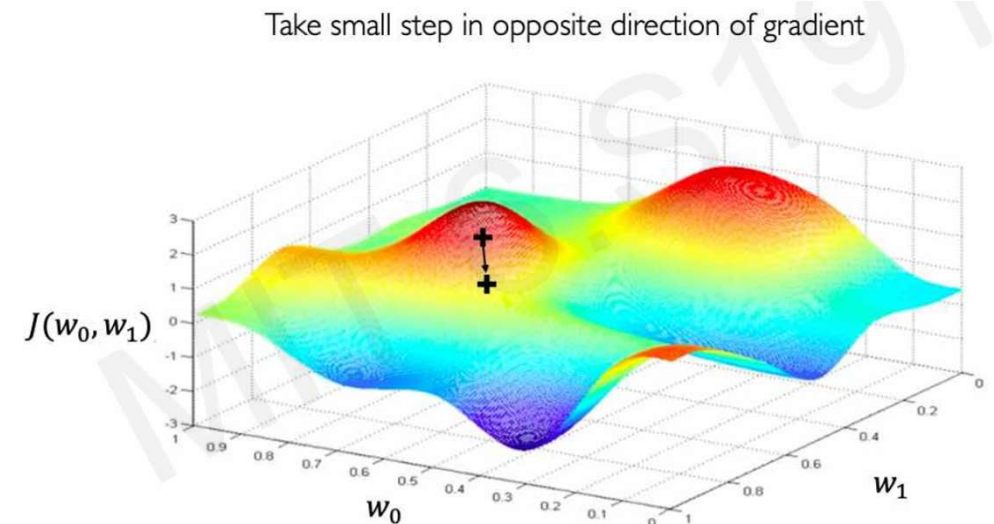
- The cost/loss function quantifies how incorrect the prediction is!
 - Function of the weights
 - Examples:
 - Binary Cross Entropy Loss – between 0 and 1
 - Mean Squared Error Loss – any real number
- **Aim of the training: Optimize/Reduce the cost function!**
- Gradient descent:
 - random initial values
 - compute gradient



<http://introtodeeplearning.com>, MIT 6.S191

How to learn/train


- The cost/loss function quantifies how incorrect the prediction is!
 - Function of the weights
 - Examples:
 - Binary Cross Entropy Loss – between 0 and 1
 - Mean Squared Error Loss – any real number
- **Aim of the training: Optimize/Reduce the cost function!**
- Gradient descent:
 - random initial values
 - compute gradient
 - move in opposite direction



<http://introtodeeplearning.com>, MIT 6.S191

How to learn/train

- The cost/loss function quantifies how incorrect the prediction is!
 - Function of the weights
 - Examples:
 - Binary Cross Entropy Loss – between 0 and 1
 - Mean Squared Error Loss – any real number
- **Aim of the training: Optimize/Reduce the cost function!**
- Gradient descent:
 - random initial values
 - compute gradient
 - move in opposite direction
 - repeat until convergence



<http://introtodeeplearning.com>, MIT 6.S191

How to learn/train - Implementation

- mean squared error loss function
 - hand calculated!
 - special case: one neuron neural network

$$J(\mathbf{W}) = \frac{1}{n} \sum_{i=1}^n \underbrace{\left(y^{(i)} - \underbrace{f(x^{(i)}; \mathbf{W})} \right)^2}_{\text{Actual} \quad \text{Predicted}}$$

<http://introtodeeplearning.com>, MIT 6.S191

```
1 import numpy as np
2 def singleNeuronGradientMeanSquaredErrorLoss(numberOfInputs, input, real_output, predicted_output, neuron):
3     """
4     Hand calculated gradient of the mean squared error loss
5     """
6     n = real_output.size
7     input = np.array(input)
8     grad = np.zeros(numberOfInputs + 1)
9     # weights gradient
10    for j in range(numberOfInputs):
11        sum = 0
12        for i in range(n):
13            sum += (real_output[i] - predicted_output[i]) * neuron.gradActivation_func(np.dot(neuron.weights, input[i,:])) * input[i,j]
14        grad[j] = -2 / n * sum
15    # bias gradient
16    for i in range(n):
17        sum += (real_output[i] - predicted_output[i]) * neuron.gradActivation_func(np.dot(neuron.weights, input[i,:]))
18    grad[-1] = -2 / n * sum
19    return grad
```

A simple neural network


- One neuron with three Inputs

- Data:

x_1	x_2	x_3	y_1
0	0	1	0
0	1	1	0
1	0	1	1
1	1	1	1
1	0	0	1

} Training data

} Testing data



<http://introtodeeplearning.com>, MIT 6.S191

A simple neural network – Implementation

```
70 numberOfTrainingIterations = 10000
71
72 train_input = np.array([[0, 0, 1],
73                          [1, 1, 1],
74                          [1, 0, 1],
75                          [0, 1, 1]])
76 train_output = np.array([0, 1, 1, 0])
77
78 test_input = np.array([[1, 0, 0]])
79 test_output = np.array([1])
80
81 learning_rate = 1.
82
83 neuron = Neuron(3, actFunc="sigmoid")
84
85 for index in range(numberOfTrainingIterations):
86     pred_output = neuron(train_input)
87     grad = singleNeuronGradientMeanSquaredErrorLoss(3, train_input, train_output, pred_output, neuron)
88     neuron.weights -= learning_rate * grad[0:-1]
89     neuron.bias -= learning_rate * grad[-1]
90
91 print("Result: ", neuron(test_input))
92 print("Weights: ", neuron.weights)
93 print("Bias: ", neuron.bias)
```

A simple neural network – Implementation

```
70 numberOfTrainingIterations = 10000
71
72 train_input = np.array([[0, 0, 1],
73                          [1, 1, 1],
74                          [1, 0, 1],
75                          [0, 1, 1]])
76 train_output = np.array([0, 1, 1, 0])
77
78 test_input = np.array([[1, 0, 0]])
79 test_output = np.array([1])
80
81 learning_rate = 1.
82
83 neuron = Neuron(3, actFunc="sigmoid")
84
85 for index in range(numberOfTrainingIterations):
86     pred_output = neuron(train_input)
87     grad = singleNeuronGradientMeanSquaredErrorLoss(3, train_input, train_output, pred_output, neuron)
88     neuron.weights -= learning_rate * grad[0:-1]
89     neuron.bias -= learning_rate * grad[-1]
90
91 print("Result: ", neuron(test_input))
92 print("Weights: ", neuron.weights)
93 print("Bias: ", neuron.bias)
```

```
Result: [0.98828371]
Weights: [ 8.88557619 -0.12700066 -2.06464961]
Bias: -4.450586368204799
```

Compare with real solution: [1]

We can look at weights and bias → supervised learning

Convolutional Neural Network

How a machine learns to see numbers

Convolution

- Mathematical definition: $(f * g)(x) := \int_{\mathbb{R}^n} f(\tau)g(x - \tau)d\tau$
 - interpretation: *weighting a function $g(x - \tau)$ with $f(\tau)$*
- Application in image processing:
 - using 2d filters to detect specific structures
 - example:

Image

1	1	0	1
0	0	1	0
0	1	0	1
0	0	0	0

+

Filter

1	0	1
0	1	0
1	0	1

→

0	0	0		
0	1	1	0	1
0	0	0	1	0
	0	1	0	1
	0	0	0	0

=

Output

1			

Convolution

- Mathematical definition: $(f * g)(x) := \int_{\mathbb{R}^n} f(\tau)g(x - \tau)d\tau$
 - interpretation: *weighting a function $g(x - \tau)$ with $f(\tau)$*
- Application in image processing:
 - using 2d filters to detect specific structures
 - example:

Image

1	1	0	1
0	0	1	0
0	1	0	1
0	0	0	0

+

Filter

1	0	1
0	1	0
1	0	1

→

0	0	0	
1	1	0	1
0	0	1	0
0	1	0	1
0	0	0	0

=

Output

1	2		

Convolution

- Mathematical definition: $(f * g)(x) := \int_{\mathbb{R}^n} f(\tau)g(x - \tau)d\tau$
 - interpretation: *weighting a function $g(x - \tau)$ with $f(\tau)$*
- Application in image processing:
 - using 2d filters to detect specific structures
 - example:

Image

1	1	0	1
0	0	1	0
0	1	0	1
0	0	0	0

+

Filter

1	0	1
0	1	0
1	0	1

→

Input

	0	0	0
1	1	0	1
0	0	1	0
0	1	0	1
0	0	0	0

=

Output

1	2	0	

Convolution

- Mathematical definition: $(f * g)(x) := \int_{\mathbb{R}^n} f(\tau)g(x - \tau)d\tau$
 - interpretation: *weighting a function $g(x - \tau)$ with $f(\tau)$*
- Application in image processing:
 - using 2d filters to detect specific structures
 - example:

Image

1	1	0	1
0	0	1	0
0	1	0	1
0	0	0	0

+

Filter

1	0	1
0	1	0
1	0	1

→

Input

1	1	0	1	0
0	0	1	0	0
0	1	0	1	—
0	0	0	0	

Output

1	2	0	2

Convolution

- Mathematical definition: $(f * g)(x) := \int_{\mathbb{R}^n} f(\tau)g(x - \tau)d\tau$
 - interpretation: *weighting a function $g(x - \tau)$ with $f(\tau)$*
- Application in image processing:
 - using 2d filters to detect specific structures
 - example:

Image

1	1	0	1
0	0	1	0
0	1	0	1
0	0	0	0

+

Filter

1	0	1
0	1	0
1	0	1

→

Input

1	1	0	1
0	0	1	0
0	1	0	1
0	0	0	0

=

Output

1	2	0	2
2	1	5	0

Convolution

- Mathematical definition: $(f * g)(x) := \int_{\mathbb{R}^n} f(\tau)g(x - \tau)d\tau$
 - interpretation: *weighting a function $g(x - \tau)$ with $f(\tau)$*
- Application in image processing:
 - using 2d filters to detect specific structures
 - example:

Image

1	1	0	1
0	0	1	0
0	1	0	1
0	0	0	0

+

Filter

1	0	1
0	1	0
1	0	1

→

Input

1	1	0	1
0	0	1	0
0	1	0	1
0	0	0	0

=

Output

1	2	0	2
2	1	5	0
0	2	0	2

Convolution

- Mathematical definition: $(f * g)(x) := \int_{\mathbb{R}^n} f(\tau)g(x - \tau)d\tau$
 - interpretation: *weighting a function $g(x - \tau)$ with $f(\tau)$*
- Application in image processing:
 - using 2d filters to detect specific structures
 - example:

Image

1	1	0	1
0	0	1	0
0	1	0	1
0	0	0	0

+

Filter

1	0	1
0	1	0
1	0	1

→

Input

1	1	0	1
0	0	1	0
0	1	0	1
0	0	0	0

=

Output


1	2	0	2
2	1	5	0
0	2	0	2
1	0	2	0

Convolutional Neural Network

- Engine: TensorFlow
- Deep neural network = a lot of neurons, in many layers
 - 2 convolutional layer with 16 and 36 neurons and one 5x5 filter each
 - the optimal filters are generated here
 - 2 fully connected layer with 128 and 10 neurons
 - the classification takes place here
- Data: MNIST database, a lot of handwritten numbers
- Credits: Magnus Erik Hvass Pedersen, [1]

[1] https://github.com/Hvass-Labs/TensorFlow-Tutorials/blob/master/02_Convolutional_Neural_Network.ipynb

Convolutional Neural Network



https://github.com/Hvass-Labs/TensorFlow-Tutorials/blob/master/02_Convolutional_Neural_Network.ipynb

Now let's look at the results!

Summary

- Neural networks...
 - ...are not that complicated
 - ...can be written from scratch
- Convolutional neural networks are quite good in computer vision. (Better than “normal” neural networks, 91 % vs 99 %)
- Supplementary material can be found here:
<https://github.com/jurvo/cp2-nn-exam>
 - Jupyter Notebooks on:
 - the simple neural network with dummy data
 - the simple neural network with titanic data (to play around)
 - the convolutional neural network